

Exercise 7 – Data Storage and File Handling

Objective

To use some of the Python 3 file handling methods, as well as the pickle and gzip modules.

Questions

1. Write a Python script to list all the unused port numbers in the `/etc/services` file between 1 and 200

Steps:

Become familiar with the input file - view it first

Write the main code to read the services file one line at a time

Use string functions or a regular expression to:

Ignore lines starting with a `#` comment character

Ignore lines that just consist of "white-space"

The `/etc/services` has several columns separated by white-space

- Use split or a regular expression to isolate the port/protocol field
- Use another split or regular expression to isolate the port number
- Don't forget to stop at port number 200!
- Note that many port numbers have > 1 entry

On Windows the file is in `'C:\WINDOWS\system32\drivers\etc\services'` or in `'C:\WINNT\system32\drivers\etc\services'`.

On OSX the file has unused ports marked as 'Unassigned'. Therefore, we have an addition requirement: ignore all lines that start with the comment delimiter `'#'`.

Many port numbers have more than one entry in the file, but you may assume they are in order.

Hints: Open the file.

Read the file line-by-line using a for loop.

Consider using a set or a dictionary to hold the port numbers.

Be careful of comparing strings and int - you will have to convert the port number to an integer.

2. Using the data in **country.txt**, construct a Python dictionary where the country name is the key and the other record details are stored in a list as the value. Store (pickle) this dictionary into a file named 'country.p'.

Notice the size of the file compared to the original, and then change the program to use gzip.

3. Now write a program which reads the pickled dictionary and displays it onto the console.

If time allows, convert your pickle to use a shelf.

If time allows...

4. This exercise uses **messier.txt**, which was used in a previous optional exercise (you do not need to have completed that exercise to do this one).

This file contains details of Messier celestial objects that are identified by a Messier number, the first field in the file.

The aim is to access the records in the file randomly, using `seek()`. First construct an index (could be in a list or a dictionary) which consists of the file position (use `tell()`) of each record. The key is the first field, the Messier number, which is prefixed M (ignore any lines that do not start with 'M').

Now prompt the user to enter a Messier number, with or without the 'M', and display the record for that celestial object.

The file **messier.txt** has a character which uses '**latin_1**', how do you cope with that?

5. You may recall an exercise from Chapter 5 Collections that timed various ways of searching for the word 'Zulu', using the program **Ex5_4.py**. The fastest technique by far was to use a dictionary.

Modify **Ex5_4.py** to use a shelf, preferably by copying the code from your dictionary implementation and modifying the copy. If you did not complete that exercise, then use the solution code in solutions/05 Collections.

You'll find that the `shelve` is considerably slower than other techniques. However, place an additional `start_timer()/end_timer()` around the `shelve` creation (including loading the data into the dictionary). This should give two sets of times for using a `shelve`, loading and searching (which is repeated in a loop).

Where is the biggest overhead? Is this a reasonable test?

Solutions

Question 1

This solution uses regular expressions and sets. A common mistake with this approach is to forget to convert the captured port number to an int, required since range returns an integer.

```
import sys
import re

if sys.platform == 'win32':
    file = r'C:\WINDOWS\system32\drivers\etc\services'
else:
    file = '/etc/services'

ports = set()

for line in open(file, 'r'):
    m = re.search(r'(\d+)/(\udp|tcp)', line)
    if m:
        port = int(m.group(1)) # Or m.groups()[0]
        if port > 200: break
        ports.add(port)

# Subtract used port numbers from full set of ports
print(set(range(1, 201)) - ports)
```

Questions 2 & 3

```
import pickle
import gzip
import shelve

# Using a compressed pickle.
country_dict = {}
for line in open('country.txt', 'r'):
    name, *row = line.split(',')
    country_dict[name] = row

outp = gzip.open('country.p', 'wb')
pickle.dump(country_dict, outp)
outp.close()

# Using a shelve.
db = shelve.open('country')
for country in country_dict.keys():
    db[country] = country_dict[country]

db.close()
db = shelve.open('country')
print(db['Belgium'])
db.close()
```

If time allows...

Question 4

```
# Construct an index.
index = []
fh_in = open('messier.txt', 'r', encoding='latin_1')

while True:
    line = fh_in.readline()
    if not line: break
    if line.startswith('M'):
        num = line[1:6].rstrip()
        index.append(fh_in.tell() - len(line))

while True:
    num = input('Enter a Messier number to exit: ')
    if num.startswith('M'):
        num = int(num[1:])
    elif num:
        num = int(num)
    else:
        num = 0

    if num < 1: break
    num -= 1

    fh_in.seek(index[num])
    print(fh_in.readline())
```

Question 5

```
import shelve

def shelve_func():
    global shelve_dict
    return shelve_dict['Zulu'] + 1

i = 0
start_timer()
shelve_dict = shelve.open('shelve_dict')

for key in words:
    shelve_dict[key] = i
    i += 1
```

```
end_timer('Shelve creation')

start_timer()

for i in range(0, LOOP_COUNT):
    line = shelve_func();

shelve_dict.close()
end_timer('Shelved dictionary')

print('Dictionary line number: ', line)
```

The timings obtained (on a test machine) were:

```
Brute_force : 1.014 seconds
Brute_force line number: 45400
Index      : 0.312 seconds
Index line number: 45400
In         : 0.296 seconds
Dictionary : 0.016 seconds
Dictionary line number: 45400
Shelve creation: 38.517 seconds
Shelved dictionary: 0.468 seconds
Dictionary line number: 45400
```

The timings are not surprising, although Python 3.2 performance does not compare well with Python 2.7, which only takes round 1.7 seconds to create the same database. The Python 2 database file sizes are considerably smaller than those used by Python 3.

After creation, caching means that we are essentially dealing with an in-memory dictionary, nevertheless, it is still slower than a conventional dictionary. The dataset is probably too small to meaningfully measure retrieval overheads. Since record sizes, and therefore the number of I/O transfers, will vary considerably between applications, a general measure with such a simple structure is not a reliable guide.